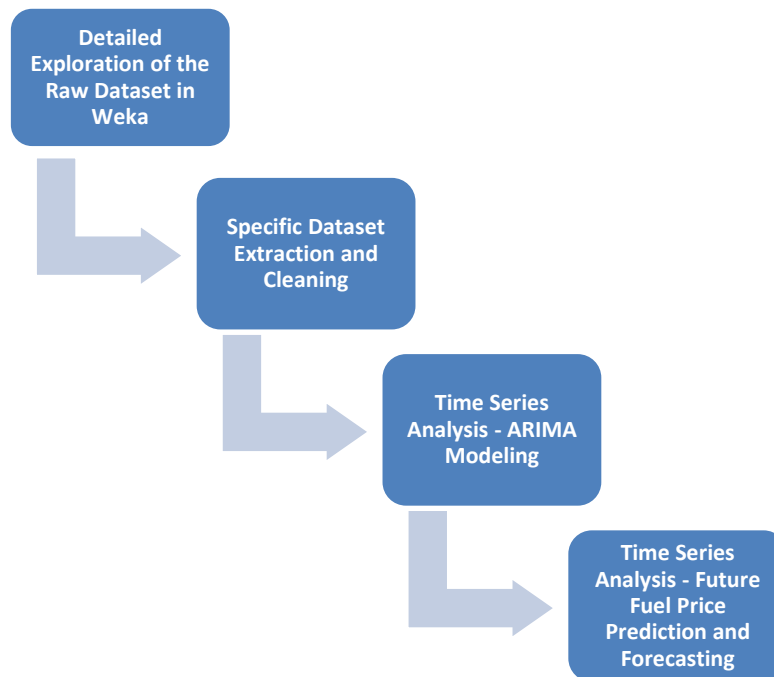


# Fuel Price Analysis and Prediction by Time Series

---

## Approach

The approach to be taken is shown in the graph below and will be described in the following subsections.



## Step 1: Detailed Exploration of the Raw Dataset in Weka

### 1. Clean and format the head (attribute names).

```
ds <- read.csv(file = "D:\\RyersonU\\CKME136 - Data Analytics - Capstone Project\\Possible  
Datasets\\Fuels price survey information\\fueltypesall.csv", header = T, stringsAsFactors = F,  
na.strings = c("", "NA"))
```

```
ds_head_clean <- as.data.frame (select(ds, -c(Type.de.carburant)))
```

```
names(ds_head_clean) <- c("DATE", "OTTAWA", "TORONTO_WEST", "TORONTO_EAST",
"WINDSOR", "LONDON", "PETERBOROUGH", "ST_CATHARINE", "SUDBURY",
"SAULT_SAINTE_MARIE", "THUNDER_BAY", "NORTH_BAY", "TIMMINS", "KENORA",
"PARRY_SOUND", "ONTARIO_AVG", "SOUTHERN_AVG", "NORTHERN_AVG", "FUEL_TYPE")
```

```
str(ds_head_clean)
```

```
> str(ds_head_clean)
'data.frame': 9422 obs. of 19 variables:
 $ DATE      : chr  "1990-01-03" "1990-01-10" "1990-01-17" "1990-01-24" ...
 $ OTTAWA    : num  55.9 55.9 55.9 55.9 55.9 55.8 55.8 55.8 55.9 56 ...
 $ TORONTO_WEST : num  49.1 47.7 53.2 53.2 51.9 50.7 49.3 48.2 54.1 53.8 ...
 $ TORONTO_EAST : num  48.7 46.8 53.2 53.5 52.6 50.7 48.4 47.1 54.4 53.8 ...
 $ WINDSOR    : num  45.2 49.7 49.6 49 48.6 48.5 48.5 48.6 48.5 48.3 ...
 $ LONDON     : num  50.1 47.6 53.7 52.1 49.1 47.8 54.7 53.6 51.7 50.8 ...
 $ PETERBOROUGH : num  0 0 0 0 0 0 0 0 0 ...
 $ ST_CATHARINE : num  0 0 0 0 0 0 0 0 0 ...
 $ SUDBURY    : num  56.4 56.4 55.8 55.7 55.6 55.6 55.6 55.6 55.6 55.6 ...
 $ SAULT_SAINTE_MARIE : num  54.8 54.9 54.9 54.9 54.8 54.8 54.8 54.8 54.8 54.8 ...
 $ THUNDER_BAY : num  56.6 56.8 56.8 56.8 56.8 56.8 56.9 57 57 57 ...
 $ NORTH_BAY   : num  55.1 55 54.4 54.3 54.2 54.2 54.1 54 54.1 54 ...
 $ TIMMINS     : num  58.1 58.2 58.2 58.2 58.1 58.1 58.1 58.1 58.1 58.1 ...
 $ KENORA      : num  0 0 0 0 0 0 0 0 0 ...
 $ PARRY_SOUND : num  0 0 0 0 0 0 0 0 0 ...
 $ ONTARIO_AVG : num  50.3 49.2 53.6 53.5 52.5 51.4 50.7 49.9 54.1 53.7 ...
 $ SOUTHERN_AVG : num  49.5 48.3 53.3 53.2 52.1 50.8 50.1 49.1 53.8 53.4 ...
 $ NORTHERN_AVG : num  56.2 56.2 56 56 55.9 55.9 55.9 55.9 55.9 55.9 ...
 $ FUEL_TYPE   : chr  "Regular Unleaded Gasoline" "Regular Unleaded Gasoline" "R
Unleaded Gasoline" ...
```

Now this dataset contains 9422 observations and 19 variables, and the attribute name is English only. Then I export this head cleaned dataset as a csv file, which will be used for the next exploration in Weka.

```
write.csv(ds_head_clean, "D:\\RyersonU\\CKME136 - Data Analytics - Capstone
Project\\Possible Datasets\\Fuels price survey information\\ds_head_clean.csv", row.names =
FALSE)
```

**2. Load the head cleaned table in Weka to look at the attribute type for each attribute.**

**Current relation**

Relation: ds\_head\_clean  
Instances: 9422

Attributes: 19  
Sum of weights: 9422

**Attributes**

All None Invert Pattern

No.	Name
1	<input checked="" type="checkbox"/> DATE
2	<input type="checkbox"/> OTTAWA
3	<input type="checkbox"/> TORONTO_WEST
4	<input type="checkbox"/> TORONTO_EAST
5	<input type="checkbox"/> WINDSOR
6	<input type="checkbox"/> LONDON
7	<input type="checkbox"/> PETERBOROUGH
8	<input type="checkbox"/> ST_CATHARINE
9	<input type="checkbox"/> SUDBURY
10	<input type="checkbox"/> SAULT_SAINTE_MARIE
11	<input type="checkbox"/> THUNDER_BAY
12	<input type="checkbox"/> NORTH_BAY
13	<input type="checkbox"/> TIMMINS
14	<input type="checkbox"/> KENORA
15	<input type="checkbox"/> PARRY_SOUND
16	<input type="checkbox"/> ONTARIO_AVG

Remove

**Selected attribute**

Name: DATE  
Missing: 0 (0%)  
Distinct: 1585  
Type: Nominal  
Unique: 0 (0%)

No.	Label	Count	Weight
1	1990-01-03	6	6.0
2	1990-01-10	6	6.0
3	1990-01-17	6	6.0
4	1990-01-24	6	6.0
5	1990-01-31	6	6.0
6	1990-02-07	6	6.0
7	1990-02-14	6	6.0
8	1990-02-21	6	6.0
9	1990-02-28	6	6.0

Class: FULE\_TYPE (Nom) Visualize All

Too many values to display.

Nominal (Qualitative) Type includes the attributes: DATE and FUEL-TYPE;  
 Numeric (Quantitative) Type includes the attributes: other 17 attributes.

### 3. Find max, min, mean and standard deviation of attributes.

By selecting each attribute name, we can get the statistics for each numeric attribute.

**Current relation**

Relation: ds\_head\_clean  
Instances: 9422

Attributes: 19  
Sum of weights: 9422

**Attributes**

All None Invert Pattern

No.	Name
1	<input type="checkbox"/> DATE
2	<input checked="" type="checkbox"/> OTTAWA
3	<input type="checkbox"/> TORONTO_WEST
4	<input type="checkbox"/> TORONTO_EAST
5	<input type="checkbox"/> WINDSOR
6	<input type="checkbox"/> LONDON
7	<input type="checkbox"/> PETERBOROUGH
8	<input type="checkbox"/> ST_CATHARINE
9	<input type="checkbox"/> SUDBURY
10	<input type="checkbox"/> SAULT_SAINTE_MARIE
11	<input type="checkbox"/> THUNDER_BAY
12	<input type="checkbox"/> NORTH_BAY
13	<input type="checkbox"/> TIMMINS
14	<input type="checkbox"/> KENORA
15	<input type="checkbox"/> PARRY_SOUND
16	<input type="checkbox"/> ONTARIO_AVG

Remove

**Selected attribute**

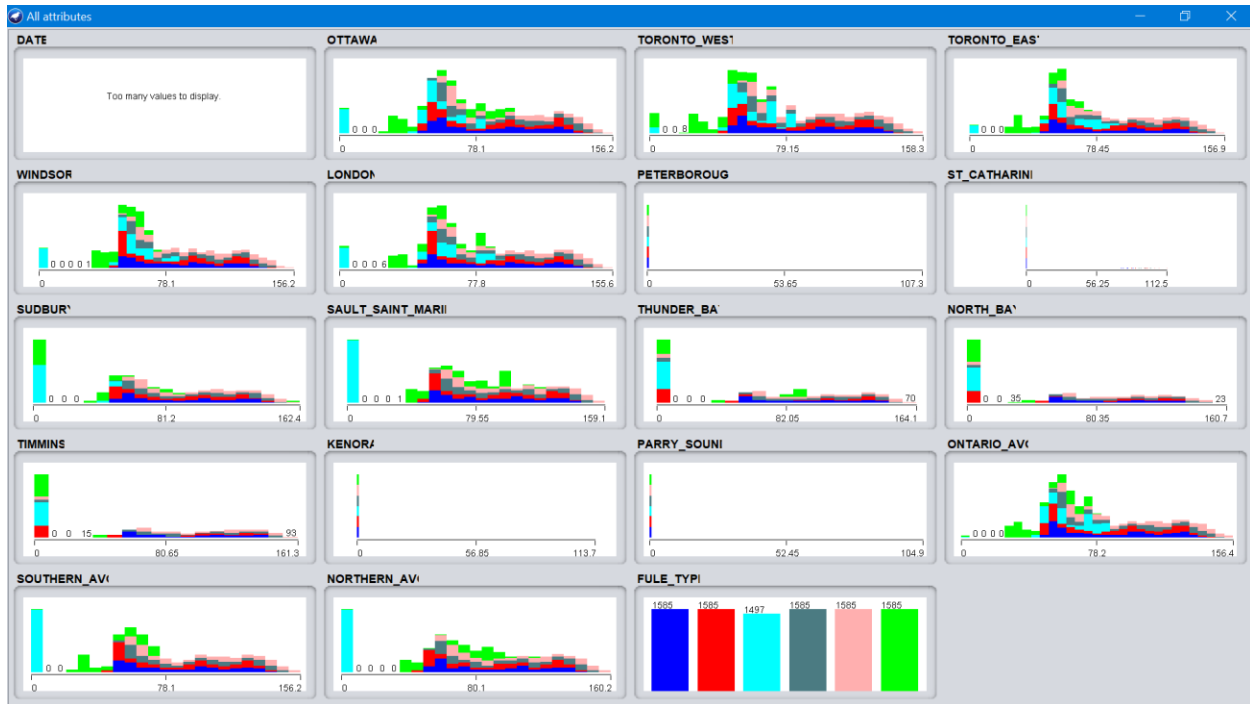
Name: OTTAWA  
Missing: 0 (0%)  
Distinct: 1113  
Type: Numeric  
Unique: 103 (1%)

Statistic	Value
Minimum	0
Maximum	156.2
Mean	76.69
StdDev	33.748

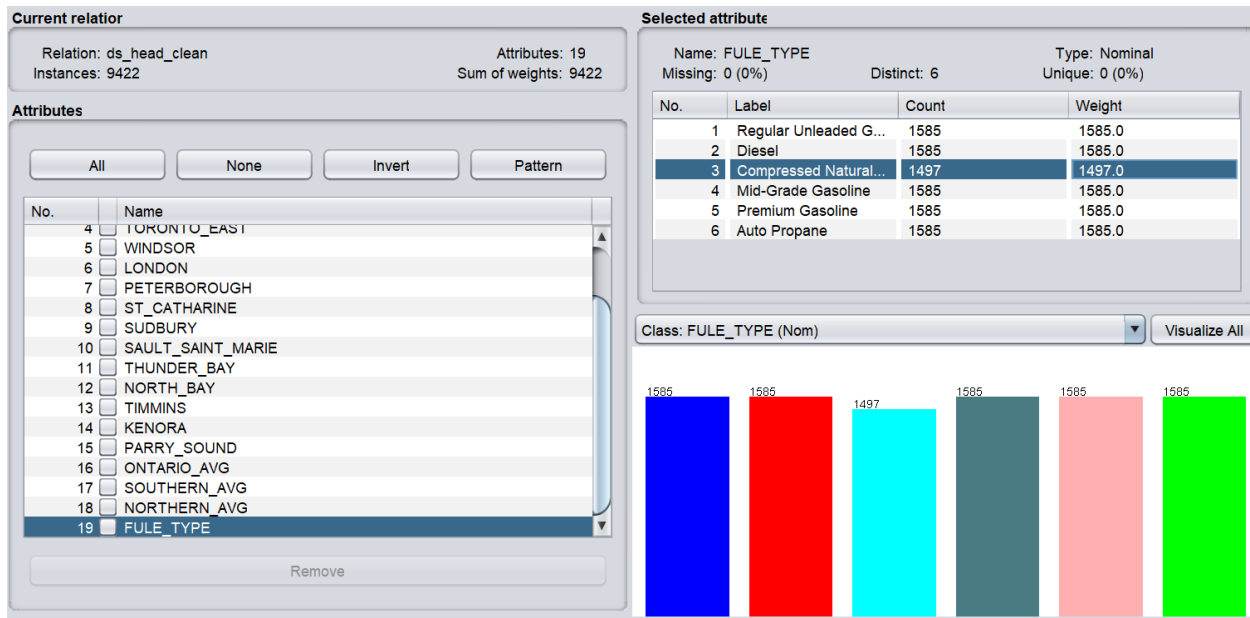
Class: FULE\_TYPE (Nom) Visualize All

The exploration indicates that all those numeric attributes do not have missing values. The contained value '0' is treated as a data value instead of missing data.

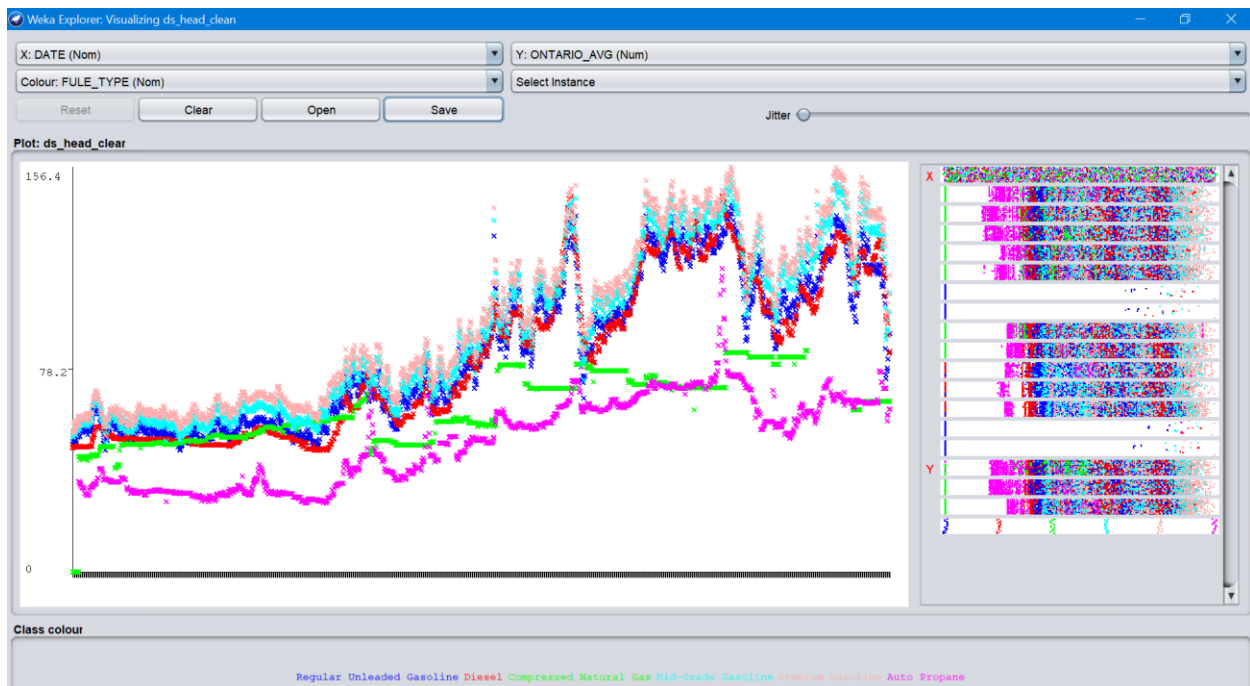
#### 4. Observe the data distribution for these attributes by visualizing all in Weka.



By observing this visualization, we can again see some cities (Markets) do not have efficient data collected. And by specifically looking at Fuel Type, we can see the type of 'Compressed Natural Gas' (blue color bar) has missing data records. This data missing will be further investigated in R in the step 2 to identify which time periods is the data missing.



## 5. Observe interesting trends in Weka Visualization.



By plotting time (Date) as X axis and Ontario Average fuel price as Y axis, we can observe a general increase trend during the period of 1991 – 2020 for all fuel types. And we can also observe some seasonal fluctuations (similar repeated pattern in a certain season) through the changes of the fuel price.

## Step 2: Specific Dataset Extraction and Cleaning

**1. Load the head-cleaned dataset (from step 1) in R and aggregate into monthly interval.**

```
RAW <- read.csv (file = "D:\\RyersonU\\CKME136 - Data Analytics - Capstone Project\\Possible  
Datasets\\Fuels price survey information\\ds_head_clean.csv", header = T, stringsAsFactors = F,  
na.strings = c("", "NA"))
```

```
RAW$YEAR_MONTH <- substr(RAW$DATE, 1, 7)
```

```
Data_Aggregated <- aggregate (x = list (RAW$OTTAWA, RAW$TORONTO_WEST,  
RAW$TORONTO_EAST, RAW$WINDSOR, RAW$LONDON, RAW$PETERBOROUGH,  
RAW$ST_CATHARINE, RAW$SUDBURY, RAW$NORTH_BAY, RAW$SAULT_SAINTE_MARIE,  
RAW$THUNDER_BAY, RAW$TIMMINS, RAW$KENORA, RAW$PARRY_SOUND), by = list  
(RAW$YEAR_MONTH, RAW$FUEL_TYPE), FUN = mean)
```

```
names (Data_Aggregated) <- c ("YEAR_MONTH", "FUEL_TYPE", "OTTAWA", "TORONTO_WEST",  
"TORONTO_EAST", "WINDSOR", "LONDON", "PETERBOROUGH", "ST_CATHARINE", "SUDBURY",  
"NORTH_BAY", "SAULT_SAINTE_MARIE", "THUNDER_BAY", "TIMMINS", "KENORA",  
"PARRY_SOUND")
```

**2. Further transfer the dataset into a new dataset with only four attributes:  
"YEAR\_MONTH", "FUEL\_TYPE", "PRICE", "CITY"**

```
OTTAWA = select (Data_Aggregated, c (YEAR_MONTH, FUEL_TYPE, OTTAWA))
```

```
OTTAWA$CITY = "OTTAWA"
```

```
names (OTTAWA) = c ("YEAR_MONTH", "FUEL_TYPE", "PRICE", "CITY")
```

```
TORONTO_WEST = select (Data_Aggregated, c (YEAR_MONTH, FUEL_TYPE, TORONTO_WEST))
```

```
TORONTO_WEST$CITY = "TORONTO_WEST"
```

*names (TORONTO\_WEST) = c ("YEAR\_MONTH", "FUEL\_TYPE", "PRICE", "CITY")*

*TORONTO\_EAST = select (Data\_Aggregated, c (YEAR\_MONTH, FUEL\_TYPE, TORONTO\_EAST))*

*TORONTO\_EAST\$CITY = "TORONTO\_EAST"*

*names (TORONTO\_EAST) = c ("YEAR\_MONTH", "FUEL\_TYPE", "PRICE", "CITY")*

*WINDSOR = select (Data\_Aggregated, c (YEAR\_MONTH, FUEL\_TYPE, WINDSOR))*

*WINDSOR\$CITY = "WINDSOR"*

*names (WINDSOR) = c ("YEAR\_MONTH", "FUEL\_TYPE", "PRICE", "CITY")*

*LONDON = select (Data\_Aggregated, c (YEAR\_MONTH, FUEL\_TYPE, LONDON))*

*LONDON\$CITY = "LONDON"*

*names (LONDON) = c ("YEAR\_MONTH", "FUEL\_TYPE", "PRICE", "CITY")*

*PETERBOROUGH = select (Data\_Aggregated, c (YEAR\_MONTH, FUEL\_TYPE, PETERBOROUGH))*

*PETERBOROUGH\$CITY = "PETERBOROUGH"*

*names (PETERBOROUGH) = c ("YEAR\_MONTH", "FUEL\_TYPE", "PRICE", "CITY")*

*ST\_CATHARINE = select (Data\_Aggregated, c (YEAR\_MONTH, FUEL\_TYPE, ST\_CATHARINE))*

*ST\_CATHARINE\$CITY = "ST\_CATHARINE"*

*names (ST\_CATHARINE) = c ("YEAR\_MONTH", "FUEL\_TYPE", "PRICE", "CITY")*

*SUDBURY = select (Data\_Aggregated, c (YEAR\_MONTH, FUEL\_TYPE, SUDBURY))*

*SUDBURY\$CITY = "SUDBURY"*

*names (SUDBURY) = c ("YEAR\_MONTH", "FUEL\_TYPE", "PRICE", "CITY")*

*NORTH\_BAY = select (Data\_Aggregated, c (YEAR\_MONTH, FUEL\_TYPE, NORTH\_BAY))*

*NORTH\_BAY\$CITY = "NORTH\_BAY"*

*names (NORTH\_BAY) = c ("YEAR\_MONTH", "FUEL\_TYPE", "PRICE", "CITY")*

*SAULT\_SAINT\_MARIE = select (Data\_Aggregated, c (YEAR\_MONTH, FUEL\_TYPE, SAULT\_SAINT\_MARIE))*

*SAULT\_SAINT\_MARIE\$CITY = "SAULT\_SAINT\_MARIE"*

```
names (SAULT_SAINT_MARIE) = c ("YEAR_MONTH", "FUEL_TYPE", "PRICE", "CITY")
```

```
THUNDER_BAY = select (Data_Aggregated, c (YEAR_MONTH, FUEL_TYPE, THUNDER_BAY))
```

```
THUNDER_BAY$CITY = "THUNDER_BAY"
```

```
names (THUNDER_BAY) = c ("YEAR_MONTH", "FUEL_TYPE", "PRICE", "CITY")
```

```
TIMMINS = select (Data_Aggregated, c (YEAR_MONTH, FUEL_TYPE, TIMMINS))
```

```
TIMMINS$CITY = "TIMMINS"
```

```
names (TIMMINS) = c ("YEAR_MONTH", "FUEL_TYPE", "PRICE", "CITY")
```

```
KENORA = select (Data_Aggregated, c (YEAR_MONTH, FUEL_TYPE, KENORA))
```

```
KENORA$CITY = "KENORA"
```

```
names (KENORA) = c ("YEAR_MONTH", "FUEL_TYPE", "PRICE", "CITY")
```

```
PARRY_SOUND = select (Data_Aggregated, c (YEAR_MONTH, FUEL_TYPE, PARRY_SOUND))
```

```
PARRY_SOUND$CITY = "PARRY_SOUND"
```

```
names (PARRY_SOUND) = c ("YEAR_MONTH", "FUEL_TYPE", "PRICE", "CITY")
```

```
City_Combined <- list (OTTAWA, TORONTO_WEST, TORONTO_EAST, WINDSOR, LONDON,  
PETERBOROUGH, ST_CATHARINE, SUDBURY, NORTH_BAY, SAULT_SAINT_MARIE, THUNDER_BAY,  
TIMMINS, KENORA, PARRY_SOUND)
```

```
City_Combined <- do.call ("rbind", City_Combined)
```

```
City_Combined <- as.data.frame (City_Combined)
```

### **3. Have a look at the new dataset: "City\_Combined"**

```
str (City_Combined)
```

```
summary (City_Combined)
```



```
> str(City_Combined)
'data.frame': 30380 obs. of 4 variables:
 $ YEAR_MONTH: chr "1990-01" "1990-02" "1990-03" "1990-04" ...
 $ FULE_TYPE : chr "Auto Propane" "Auto Propane" "Auto Propane" "Auto Propane" ...
 $ PRICE : num 0 0 0 30.8 30.6 ...
 $ CITY : chr "OTTAWA" "OTTAWA" "OTTAWA" "OTTAWA" ...
> summary(City_Combined)
YEAR_MONTH      FULE_TYPE      PRICE      CITY
Length:30380    Length:30380    Min. : 0.00    Length:30380
Class :character Class :character 1st Qu.: 0.00    Class :character
Mode :character  Mode :character Median : 51.90    Mode :character
                        Mean : 49.61
                        3rd Qu.: 86.90
                        Max. :162.97
```

```
First_10_Row <- head(City_Combined, n=10)
```

	YEAR_MONTH	FULE_TYPE	PRICE	CITY
1	1990-01	Auto Propane	0.000	OTTAWA
2	1990-02	Auto Propane	0.000	OTTAWA
3	1990-03	Auto Propane	0.000	OTTAWA
4	1990-04	Auto Propane	30.800	OTTAWA
5	1990-05	Auto Propane	30.560	OTTAWA
6	1990-06	Auto Propane	30.350	OTTAWA
7	1990-07	Auto Propane	28.625	OTTAWA
8	1990-08	Auto Propane	31.240	OTTAWA
9	1990-09	Auto Propane	33.400	OTTAWA
10	1990-10	Auto Propane	37.140	OTTAWA

```
Last_10_Row <- tail(City_Combined, n=10)
```

	YEAR_MONTH	FULE_TYPE	PRICE	CITY
30371	2019-08	Regular Unleaded Gasoline	0.000	PARRY_SOUND
30372	2019-09	Regular Unleaded Gasoline	0.000	PARRY_SOUND
30373	2019-10	Regular Unleaded Gasoline	0.000	PARRY_SOUND
30374	2019-11	Regular Unleaded Gasoline	0.000	PARRY_SOUND
30375	2019-12	Regular Unleaded Gasoline	0.000	PARRY_SOUND
30376	2020-01	Regular Unleaded Gasoline	0.000	PARRY_SOUND
30377	2020-02	Regular Unleaded Gasoline	0.000	PARRY_SOUND
30378	2020-03	Regular Unleaded Gasoline	0.000	PARRY_SOUND
30379	2020-04	Regular Unleaded Gasoline	74.375	PARRY_SOUND
30380	2020-05	Regular Unleaded Gasoline	81.550	PARRY_SOUND

## 4. Identify and Remove Missing Values

Based on the general exploration in Step 1, I found there are missing values under the fuel type of Compressed Natural Gas. Now I am investigating those missing values in R.

### 1) To confirm the missing value is for Compressed Natural Gas

```
table (City_Combined$FUEL_TYPE)
```

```
> table (City_Combined$FULE_TYPE)
```

Auto Propane	Compressed Natural Gas	Diesel
5110	4830	5110
Mid-Grade Gasoline	Premium Gasoline Regular Unleaded Gasoline	
5110	5110	5110

### 2) Introduce a column of YEAR to help checking which year(s) contains the missing value

```
City_Combined$YEAR <- substr(City_Combined$YEAR_MONTH, 1, 4)
```

```
table (City_Combined$YEAR)
```

```
> table(City_Combined$YEAR)
```

1990	1991	1992	1993	1994	1995	1996	1997	1998	1999	2000	2001	2002	2003	2004	2005
1008	1008	1008	1008	1008	1008	1008	1008	1008	1008	1008	1008	1008	1008	1008	1008
2006	2007	2008	2009	2010	2011	2012	2013	2014	2015	2016	2017	2018	2019	2020	
1008	1008	1008	1008	1008	1008	1008	1008	1008	1008	1008	896	840	1008	420	

We can see the missing value is in 2017 and 2018. We can ignore 2020 because the original data collected is only up to May 2020.

### 3) Final confirms for missing value

```
CNG <- subset (City_Combined, FUEL_TYPE == "Compressed Natural Gas")
```

```
table (CNG$YEAR)
```

```
> CNG <- subset(City_Combined, FULE_TYPE == "Compressed Natural Gas")
> table(CNG$YEAR)
```

1990	1991	1992	1993	1994	1995	1996	1997	1998	1999	2000	2001	2002	2003	2004	2005
168	168	168	168	168	168	168	168	168	168	168	168	168	168	168	168
2006	2007	2008	2009	2010	2011	2012	2013	2014	2015	2016	2017	2019	2020		
168	168	168	168	168	168	168	168	168	168	168	56	168	70		

Therefore, the missing value is caused by that the Compressed Natural Gas (CNG) only has partial data in year of 2017 and does not have any data in year of 2018. So, I decided to remove the CNG fuel type from my dataset.

```
City_Combined_NoCNG <- subset (City_Combined, FUEL_TYPE != "Compressed Natural Gas")
table (City_Combined_NoCNG$FUEL_TYPE)
```

```
> table(City_Combined_NoCNG$FUEL_TYPE)
```

Auto Propane	Diesel	Mid-Grade Gasoline
5110	5110	5110
Premium Gasoline	Regular Unleaded Gasoline	
5110	5110	

## 5. Generate Subsets Based on different types of Fuel and Identify 0 values

### 1) Generate 5 subsets and Plotting them

```
AP <- subset (City_Combined_NoCNG, FUEL_TYPE == "Auto Propane")
```

```
DSL <- subset (City_Combined_NoCNG, FUEL_TYPE == "Diesel")
```

```
MGG <- subset (City_Combined_NoCNG, FUEL_TYPE == "Mid-Grade Gasoline")
```

```
RUG <- subset (City_Combined_NoCNG, FUEL_TYPE == "Regular Unleaded Gasoline")
```

```
PRG <- subset (City_Combined_NoCNG, FUEL_TYPE == "Premium Gasoline")
```

```
Plot_AP <- ggplot(AP, aes(x = YEAR, y = PRICE, group = CITY)) +
```

```
geom_line(aes(color=CITY)) + geom_point(aes(color=CITY))
```

```
Plot_DSL <- ggplot(DSL, aes(x = YEAR, y = PRICE, group = CITY)) +
```

```
geom_line(aes(color=CITY)) + geom_point(aes(color=CITY))
```

```
Plot_MGG <- ggplot(MGG, aes(x = YEAR, y = PRICE, group = CITY)) +
```

```
geom_line(aes(color=CITY)) + geom_point(aes(color=CITY))
```


```
Plot_RUG <- ggplot(RUG, aes(x = YEAR, y = PRICE, group = CITY)) +
```

```
geom_line(aes(color=CITY)) + geom_point(aes(color=CITY))
```

```
Plot_PRG <- ggplot(PRG, aes(x = YEAR, y = PRICE, group = CITY)) +
```

```
geom_line(aes(color=CITY)) + geom_point(aes(color=CITY))
```

The below is the plot of Auto Propane and we can see there are some cities have 0 value.



## 2) Identify 0 values for each Fuel Type

# AP Type (Auto Propane) #

*table (AP\$PRICE==0, AP\$CITY)*

```
> table(AP$PRICE==0, AP$CITY)
```

	KENORA	LONDON	NORTH_BAY	OTTAWA	PARRY_SOUND	PETERBOROUGH
FALSE	2	359	45	362	0	0
TRUE	363	6	320	3	365	365

	SAULT_SAINT_MARIE	ST_CATHARINE	SUDBURY	THUNDER_BAY	TIMMINS	TORONTO_EAST
FALSE		362	0	172	181	46
TRUE		3	365	193	184	319

	TORONTO_WEST	WINDSOR
FALSE	313	361
TRUE	52	4

“TRUE” indicates the count of 0 values contained in each city. We can see that all cities have 0 values for Auto Propane price and some of them contains only 0 rather than other valid values.

# DSL Type (Diesel) #

*table (DSL\$PRICE==0, DSL\$CITY)*

```
> table(DSL$PRICE==0, DSL$CITY)
```

	KENORA	LONDON	NORTH_BAY	OTTAWA	PARRY_SOUND	PETERBOROUGH
FALSE	2	365	200	365	2	2
TRUE	363	0	165	0	363	363

	SAULT_SAINT_MARIE	ST_CATHARINE	SUDBURY	THUNDER_BAY	TIMMINS	TORONTO_EAST
FALSE		365	2	365	200	200
TRUE		0	363	0	165	165

	TORONTO_WEST	WINDSOR
FALSE	365	365
TRUE	0	0

These cities do not contain 0 values for Diesel price: LONDON, OTTAWA, SAULT\_SAINT\_MARIE, SUDBURY, TORONTO\_EAST, TORONTO\_WEST, WINDSOR

# MGG Type (Mid-Grade Gasoline) #

*table (MGG\$PRICE==0, MGG\$CITY)*

```
> table(MGG$PRICE==0, MGG$CITY)
```

	KENORA	LONDON	NORTH_BAY	OTTAWA	PARRY_SOUND	PETERBOROUGH
FALSE	2	365	321	365	2	2
TRUE	363	0	44	0	363	363

	SAULT_SAINT_MARIE	ST_CATHARINE	SUDBURY	THUNDER_BAY	TIMMINS	TORONTO_EAST
FALSE		365	2	365	321	322
TRUE		0	363	0	44	43

	TORONTO_WEST	WINDSOR
FALSE	365	365
TRUE	0	0

These cities do not contain 0 values for Mid-Grade Gasoline price: LONDON, OTTAWA, SAULT\_SAINT\_MARIE, SUDBURY, TORONTO\_EAST, TORONTO\_WEST, WINDSOR.

# RUG Type (Regular Unleaded Gasoline) #

*table (RUG\$PRICE==0, RUG\$CITY)*

```
> table(RUG$PRICE==0, RUG$CITY)
```

	KENORA	LONDON	NORTH_BAY	OTTAWA	PARRY_SOUND	PETERBOROUGH
FALSE	2	365	364	365	2	2
TRUE	363	0	1	0	363	363

	SAULT_SAINTE_MARIE	ST_CATHARINE	SUDBURY	THUNDER_BAY	TIMMINS	TORONTO_EAST
FALSE	365	2	365	364	364	365
TRUE	0	363	0	1	1	0

	TORONTO_WEST	WINDSOR
FALSE	365	365
TRUE	0	0

These cities do not contain 0 values for Regular Unleaded Gasoline price: LONDON, OTTAWA, SAULT\_SAINTE\_MARIE, SUDBURY, TORONTO\_EAST, TORONTO\_WEST, WINDSOR.

# PRG Type (Premium Gasoline) #

```
table (PRG$PRICE==0, PRG$CITY)
```

```
> table(PRGS$PRICE==0, PRGS$CITY)
```

	KENORA	LONDON	NORTH_BAY	OTTAWA	PARRY_SOUND	PETERBOROUGH
FALSE	2	365	321	365	2	2
TRUE	363	0	44	0	363	363

	SAULT_SAINTE_MARIE	ST_CATHARINE	SUDBURY	THUNDER_BAY	TIMMINS	TORONTO_EAST
FALSE	365	2	365	321	322	365
TRUE	0	363	0	44	43	0

	TORONTO_WEST	WINDSOR
FALSE	365	365
TRUE	0	0

These cities do not contain 0 values for Premium Gasoline price: LONDON, OTTAWA, SAULT\_SAINTE\_MARIE, SUDBURY, TORONTO\_EAST, TORONTO\_WEST, WINDSOR

### 3) Cleaning 0 values for each Fuel Type

Because the 0 value is treated as invalid value, I will clean those 0 values by removing those cities who contain 0 values and keep only those cities who does not contain 0 values for future my analysis. Therefore, I will drop the Auto Propane (AP) fuel type, and keep only Diesel (DSL), Mid-Grade Gasoline (MGG), Regular Unleaded Gasoline (RUG), and Premium Gasoline (PRG) fuel types. I will also keep only the cities of LONDON,

OTTAWA, SAULT\_SAINT\_MARIE, SUDBURY, TORONTO\_EAST, TORONTO\_WEST, and WINDSOR for my future analysis.

```
DSL_Clean <- subset (DSL, grepl  
( 'LONDON|OTTAWA|SAULT_SAINT_MARIE|SUDBURY|TORONTO|WINDSOR', DSL$CITY))  
MGG_Clean <- subset (MGG, grepl  
( 'LONDON|OTTAWA|SAULT_SAINT_MARIE|SUDBURY|TORONTO|WINDSOR',  
MGG$CITY))  
RUG_Clean <- subset (RUG, grepl  
( 'LONDON|OTTAWA|SAULT_SAINT_MARIE|SUDBURY|TORONTO|WINDSOR',  
RUG$CITY))  
PRG_Clean <- subset (PRG, grepl  
( 'LONDON|OTTAWA|SAULT_SAINT_MARIE|SUDBURY|TORONTO|WINDSOR',  
PRG$CITY))
```

## 6. Get Final Subsets Ready for Time Series Analysis

To get final subsets ready, I will drop the previously introduced column of YEAR which is used for missing value investigation. So, in each subset, there will be four attributes: YEAR\_MONTH, FUEL\_TYPE, PRICE, CITY.

```
DSL_Ready <- select (DSL_Clean, -c(YEAR))  
MGG_Ready <- select (MGG_Clean, -c(YEAR))  
RUG_Ready <- select (RUG_Clean, -c(YEAR))  
PRG_Ready <- select (PRG_Clean, -c(YEAR))
```

	YEAR_MONTH	FUEL_TYPE	PRICE	CITY
--	------------	-----------	-------	------

Therefore, the time series analysis in my project will be an analysis of the prices for four different fuel types (Diesel, Mid-Grade Gasoline, Regular Unleaded Gasoline, Premium Gasoline) based on seven different Ontario marketplaces (Ottawa, Toronto West, Toronto East, Windsor, London, Sudbury, Sault Saint Marie).

### Step 3: Time Series Analysis - ARIMA Modeling

ARIMA stands for Autoregressive Integrated Moving Average. Auto Regressive (AR) refers to the lags of the differenced series, Moving Average (MA) refers to the lags of errors, and (I) refers to the number of differences used to make the time series stationary.

The Assumptions of ARIMA model include:

- Data should be stationary: the series does not depend on the time when it is captured. White noise series and series with cyclic behavior can be considered as stationary series.
- Data should be univariate: ARIMA works on a single variable and Autoregression is regression with the past values.

My ARIMA modeling will include three sub-steps: Time Series Exportation, Model Fitting, and Measures Evaluation.

R package will be used are as following:

```
library(fUnitRoots)
```

```
library(lmtest)
```

```
library(forecast)
```

```
library(FitAR)
```

#### 1. Time Series Exportation

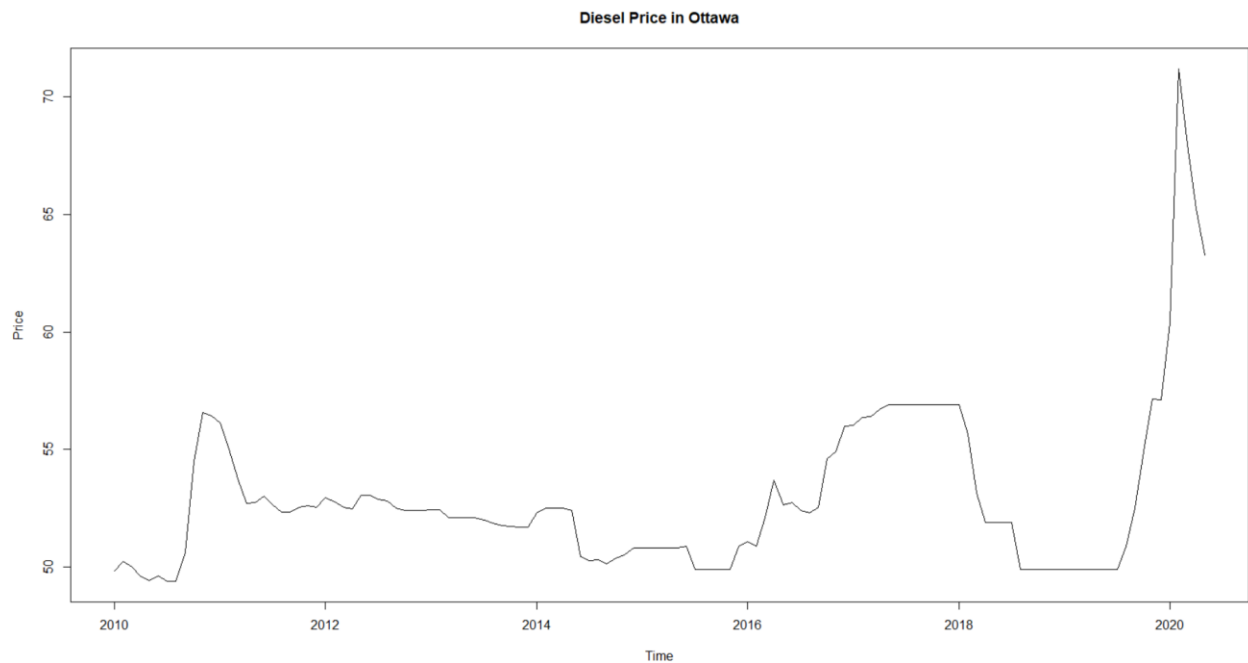
##### 1) Data Converting to Time Series

To convert data, I did a further subset of the DSL\_Ready dataset to focus Diesel price only in Ottawa marketplace first. In `ts()` function, I select the "PRICE" attribute as the univariate, and I shorten the time period to Jan 2010 – May 2020 in order to achieve a better view for the following graphing. Frequency is set to 12 because of monthly interval data.

```
DSL_Ottawa <- subset(DSL_Ready, CITY == "OTTAWA")
```



```
ts_DSL_Ottawa = ts (DSL_Ottawa$PRICE, start = c (2010,1), end = c (2020,5),
frequency=12)
plot (ts_DSL_Ottawa, main="Diesel Price in Ottawa", ylab="Price")
```




I can see from the graph that the data follows an overall upward trend. Next I will conduct the Autocorrelation analysis to examine serial dependence. And then I will analyze the components of the time series. After that, I will find the non-stationarity and seasonality in the data.

## 2) Autocorrelation Analysis – ACF Plotting

To calculate autocorrelation, I applied the `acf()` function and achieved the graph below.

```
acf (ts_DSL_Ottawa, lag.max=34)
```




The autocorrelation at lag 0 is included by default which always takes the value 1 because it represents the correlation between the data and themselves. From this above graph, I can see that the autocorrelation (ACF) is decreasing as the Lag increases. This means that no linear association between observations separated by larger lags.

### 3) Components Analysis

```
TSC_ts_DSL_Ottawa <- decompose(ts_DSL_Ottawa)
```

```
plot(TSC_ts_DSL_Ottawa)
```

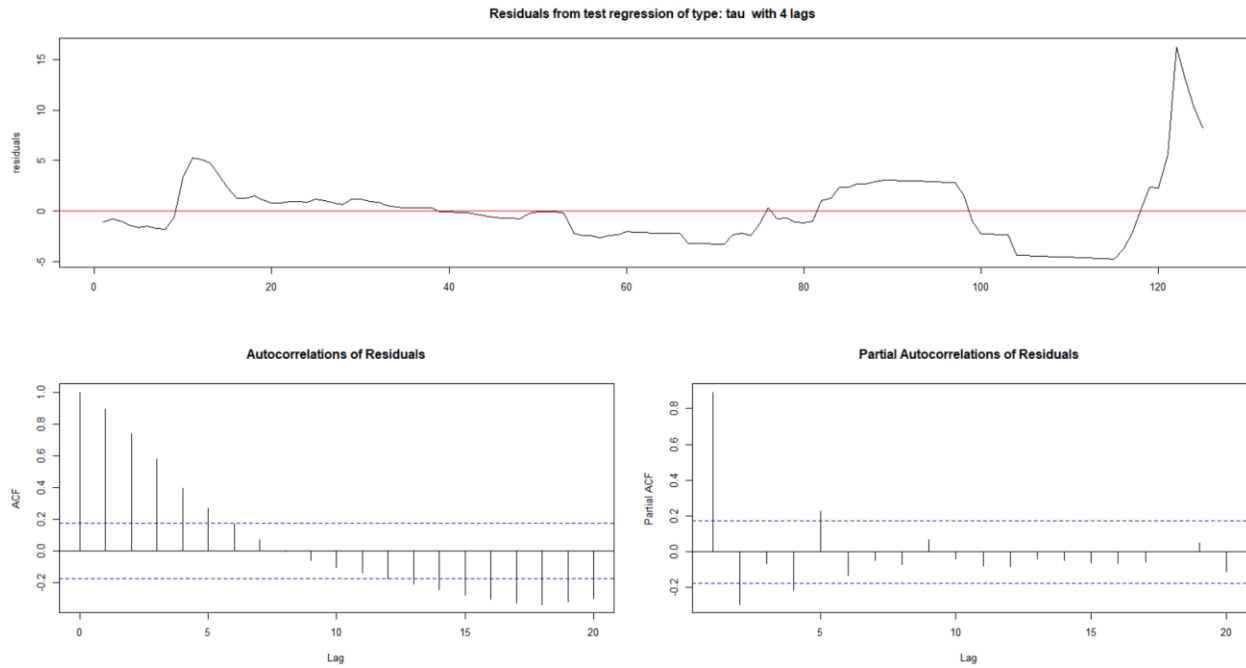


From above ACP plot, I get four components: Observed (the actual data plot), Trend (the overall upward or downward movement), Seasonal (yearly/monthly pattern of the data), Random (unexplainable part of the data).

#### 4) Determine Stationarity of data and Remove Non-Stationarity


The process of determining stationarity includes: Difference the data (compute the differences between consecutive observations), Log or Square Root the series data (stabilize non-constant variance), Fit curve to the data and Model residuals from the fit (if the data contains a trend), Unit Root Test (find first difference or regression used on the trending data to make it stationary. In Kwiatkowski-Phillips-Schmidt-Shin (KPSS) test, small p-values suggest that differencing is required.

```
urkpssTest(ts_DSL_Ottawa, type = c("tau"), lags = c("short"), use.lag = NULL, doplot = TRUE)
```



To remove non-stationarity, I applied the below code and achieved the graph after non-stationarity removal.


```
ST_ts_DSL_Ottawa <- diff(ts_DSL_Ottawa, differences=1)
plot(ST_ts_DSL_Ottawa)
```



### 5) Determine Seasonality of data and Remove Seasonality


To determine and remove seasonality from the data, I first subtract the seasonal component from the original series. Then I difference it to make it stationary. The below graph is the seasonality after the subtraction.

```
SE_ts_DSL_Ottawa_adjusted <- ts_DSL_Ottawa - TSC_ts_DSL_Ottawa$seasonal  
plot (SE_ts_DSL_Ottawa_adjusted)
```



And this below is the graph after removing seasonality and making stationary.

```
SE_ts_DSL_Ottawa <- diff (SE_ts_DSL_Ottawa_adjusted, differences=1)  
plot (SE_ts_DSL_Ottawa)
```



## 2. Model Fitting

### 1) Examine p and q values

We will need the three variables:  $p$ ,  $d$ ,  $q$  to determine the order of the model to be fitted to the data. The three variables are non-negative integers that represent the order of the autoregressive, integrated, and moving average parts of the model respectively.

I applied `pacf ()` function and `acf ()` function to examine appropriate values of  $p$  and  $q$ . The `pacf ()` is autocorrelation function at lag  $k$ . The function describes the correlation between those data points which are exactly  $k$  steps apart. It helps to identify the number of AR coefficients ( $p$ -value) in an ARIMA model. The `acf ()` function is to define values of  $p$  and  $q$  by looking at the shape of the graph. Combines with the below table, we can determine which type of the model to select and the values of  $p$ ,  $d$  and  $q$ .


Shape	Indicated Model
Exponential series decaying to 0	Auto Regressive (AR) model. <code>pacf()</code> function to be used to identify the order of the model
Alternative positive and negative spikes, decaying to 0	Auto Regressive (AR) model. <code>pacf()</code> function to be used to identify the order of the model
One or more spikes in series, rest all are 0	Moving Average(MA) model, identify order where plot becomes 0
After a few lags overall a decaying series	Mixed AR & MA model
Total series is 0 or nearly 0	Data is random
Half values at fixed intervals	We need to include seasonal AR term
Visible spikes, no decay to 0	Series is not stationary

The R code to run `pacf ()` function and `acf ()` function for this dataset is shown below associated with the result graph. The purpose of `par ()` function is to make the two graphs display in one plot view.

```
par (mfrow=c (2,1))
```

```
pacf (SE_ts_DSL_Ottawa, lag.max=34)
```

```
acf (SE_ts_DSL_Ottawa, lag.max=34)
```



## 2) Fit the Model

To fit the model, I applied the `arima ()` function with decided parameter values.

```
fitARIMA_ts_DSL_Ottawa <- arima (ts_DSL_Ottawa, order=c (1,1,1), seasonal = list  
(order = c (1,0,0), period = 12), method="ML")
```

The “order” parameter indicates the non-seasonal part of the ARIMA model and the order of (p,d,q) is the AR, the degree of difference, and the MA. The “seasonal” parameter indicates the seasonal part of the ARIMA model. The period is the frequency which is 12 (monthly) in this case. And this “seasonal” parameter requires a list of components for order and period, but it will give a numeric vector of length 3 and then turn into a suitable list with the specification as the “order”. The “method” parameter indicates the fitting method, including ML (Maximum Likelihood) and CSS (Minimize Conditional Sum-of-Squares). The default is CSS. And the ML will maximize the log-likelihood for the given (p,d,q) values to maximize the probability of obtaining the observed data.

The output of the ARIMA model can be seen by just typing the model name:

```
fitARIMA_ts_DSL_Ottawa
```

```
> fitARIMA_ts_DSL_Ottawa
```

```
Call:
```

```
arima(x = ts_DSL_Ottawa, order = c(1, 1, 1), seasonal = list(order = c(1, 0,  
0), period = 12), method = "ML")
```

```
Coefficients:
```

```
      ar1      ma1      sar1  
-0.1876  0.4528  0.0215  
s.e.    0.2739  0.2498  0.1762
```

```
sigma^2 estimated as 1.666:  log likelihood = -207.62,  aic = 423.25
```



From the output, we can check the fitted coefficients with standard error (s.e.) for each coefficient. By checking the coefficients, we can exclude the insignificant ones. We can use a function `confint()` for this purpose. We can also use `coeftest()` function for Z test of the coefficients to check p value and significant level.

```
confint(fitARIMA_ts_DSL_Ottawa)
```

```
coeftest(fitARIMA_ts_DSL_Ottawa)
```

```
> confint(fitARIMA_ts_DSL_Ottawa)
              2.5 %      97.5 %
ar1  -0.72444068  0.3492247
ma1   -0.03673989  0.9423868
sar1  -0.32394342  0.3669358
> coeftest(fitARIMA_ts_DSL_Ottawa)

z test of coefficients:

      Estimate Std. Error z value Pr(>|z|)
ar1   -0.187608   0.273899  -0.6850  0.49337
ma1    0.452823   0.249782   1.8129  0.06985 .
sar1    0.021496   0.176248   0.1220  0.90293
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The model fitting can be a repetitive process because we need to adjust different (p,d,q) values to find the most efficient and optimized model. Additionally, we can also apply the `auto.arima()` function to automate the ARIMA modeling, which will automatically help us decide the best (p,d,q) values. The `auto.arima()` function in R uses a combination of the Unit Root Test and the minimization of AIC and MLE to obtain the ARIMA model.

For example, I applied the `auto.arima()` function in R and achieve the following result:

```
auto.arima(ts_DSL_Ottawa, trace=TRUE)
```

```

> auto.arima(ts_DSL_Ottawa, trace=TRUE)

ARIMA(2,1,2)(1,0,1)[12] with drift : Inf
ARIMA(0,1,0) with drift : 427.1068
ARIMA(1,1,0)(1,0,0)[12] with drift : 425.1892
ARIMA(0,1,1)(0,0,1)[12] with drift : 423.5177
ARIMA(0,1,0) : 425.8586
ARIMA(0,1,1) with drift : 421.4081
ARIMA(0,1,1)(1,0,0)[12] with drift : 423.5208
ARIMA(0,1,1)(1,0,1)[12] with drift : 425.349
ARIMA(1,1,1) with drift : 423.048
ARIMA(0,1,2) with drift : 422.9383
ARIMA(1,1,0) with drift : 423.0807
ARIMA(1,1,2) with drift : 425.0865
ARIMA(0,1,1) : 419.8122
ARIMA(0,1,1)(1,0,0)[12] : 421.8992
ARIMA(0,1,1)(0,0,1)[12] : 421.8975
ARIMA(0,1,1)(1,0,1)[12] : 423.7008
ARIMA(1,1,1) : 421.464
ARIMA(0,1,2) : 421.3923
ARIMA(1,1,0) : 421.4619
ARIMA(1,1,2) : 423.5219

Best model: ARIMA(0,1,1)

Series: ts_DSL_Ottawa
ARIMA(0,1,1)

Coefficients:
      ma1
      0.2807
s.e.  0.0926

sigma^2 estimated as 1.686: log likelihood=-207.86
AIC=419.71 AICc=419.81 BIC=425.35

```

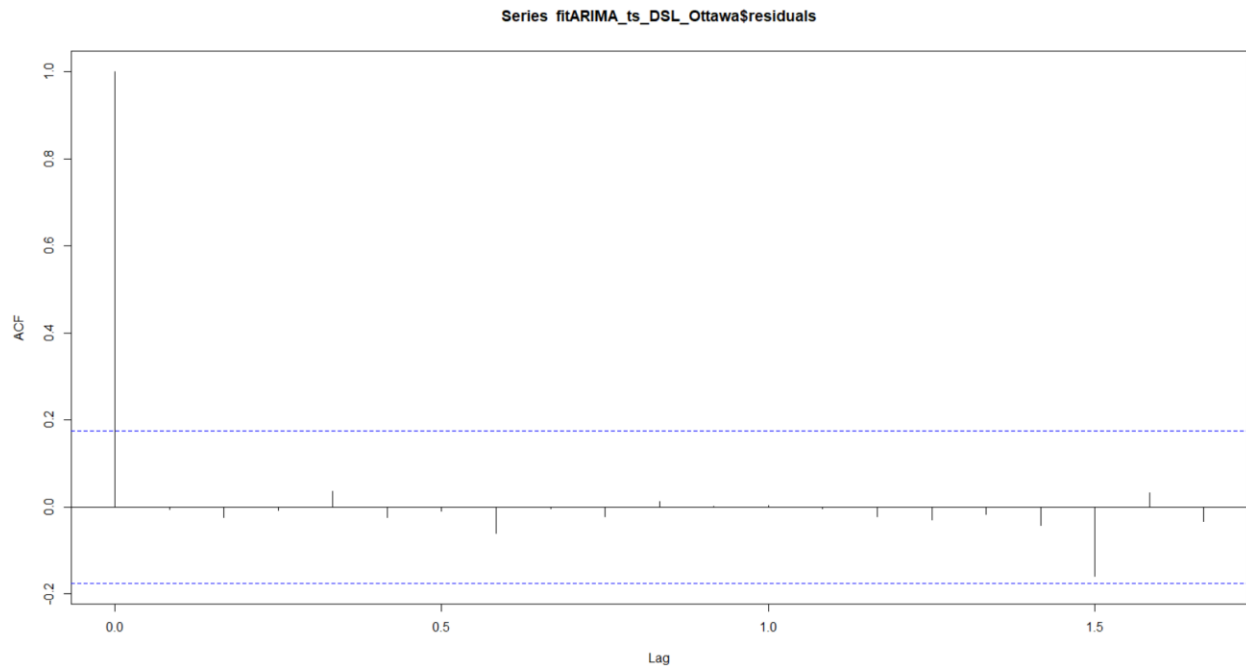
The result indicates that the automation selected ARIMA (0,1,1) as the best model.

### 3. Measures Evaluation

In measures evaluation, I use the `acf()` function to plot the pattern in the residuals of the model, which is to check if the pattern would be like white noise. Because we will only do forecasting when residuals are like white noise, otherwise, we will need to modify the selected model. I will also do a portmanteau test for the residuals of the model. The portmanteau test is to test the independence at all lags up to the one specified, which is an overall randomness test based on several lags. Because the portmanteau test is on the residuals, its hypothesis tested is also for the residuals of the selected ARIMA model and does not have autocorrelation.

## 1) ACF Plotting

*acf (fitARIMA\_ts\_DSL\_Ottawa\$residuals)*

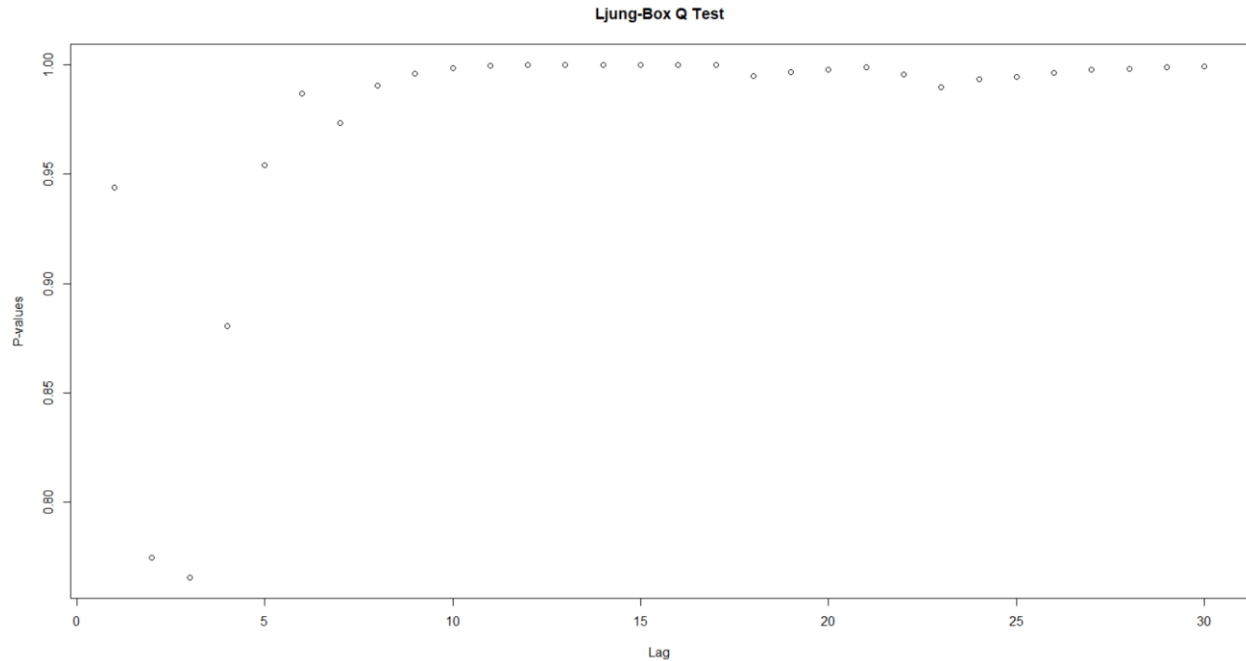


The ACF plotting shows that the residuals do not have significant autocorrelations.


## 2) Portmanteau Testing - Ljung–Box Q test

For portmanteau testing, I use Ljung–Box Q test.

```
boxresult <- LjungBoxTest (fitARIMA_ts_DSL_Ottawa$residuals,k=2,StartLag=1)
plot(boxresult[,3],main= "Ljung-Box Q Test", ylab= "P-values", xlab= "Lag")
qqnorm (fitARIMA_ts_DSL_Ottawa$residuals)
qqline (fitARIMA_ts_DSL_Ottawa$residuals)
```



From the result graph, we can see that the p-values for the Ljung–Box Q test are all greater than 0.05, which means non-significance.



And the above Normal graph indicates the values are normal because they are almost aligning on a line.

Therefore, based on all the graphs generated in this measure evaluation process, it is indicating that there is no pattern in the residuals. So, I can proceed to the Step 4 of the Approach section to conduct the forecasting.

## Step 4: Time Series Analysis - Future Fuel Price Prediction and Forecasting

The selected ARIMA model can be used as a predictive model for the forecasting of future values of the time series. To run prediction in R, there are two functions I found, one is the `predict ()` function, another is the `forecast.Arima ()` function.

### 1. The `predict ()` function

In this function, the “n.ahead” argument represents how many time steps ahead to predict. In this case, I used 12 as a one-year ahead prediction. So, I will archive the predicted value from June 2020 to May 2021.

```
predict (fitARIMA_ts_DSL_Ottawa, n.ahead = 12)
```

```
> predict(fitARIMA_ts_DSL_Ottawa,n.ahead = 12)
$pred
      Jan      Feb      Mar      Apr      May      Jun      Jul      Aug      Sep      Oct      Nov      Dec
2020 62.78422 62.78422 62.80507 62.83842 62.88846 62.93536 62.93432
2021 63.00312 63.22827 63.16114 63.10422 63.06274

$se
      Jan      Feb      Mar      Apr      May      Jun      Jul      Aug      Sep      Oct      Nov      Dec
2020 1.292936 2.100865 2.674992 3.146035 3.555205 3.921917 4.257157
2021 4.567859 4.858733 5.133150 5.393624 5.642086
```

### 2. The `forecast.Arima ()` function

In this function, the “h” argument represents how many time steps ahead. In this case, I keep the same time ahead which is the 12-month forecasting. The result contains both 80% prediction interval and 95% prediction interval.

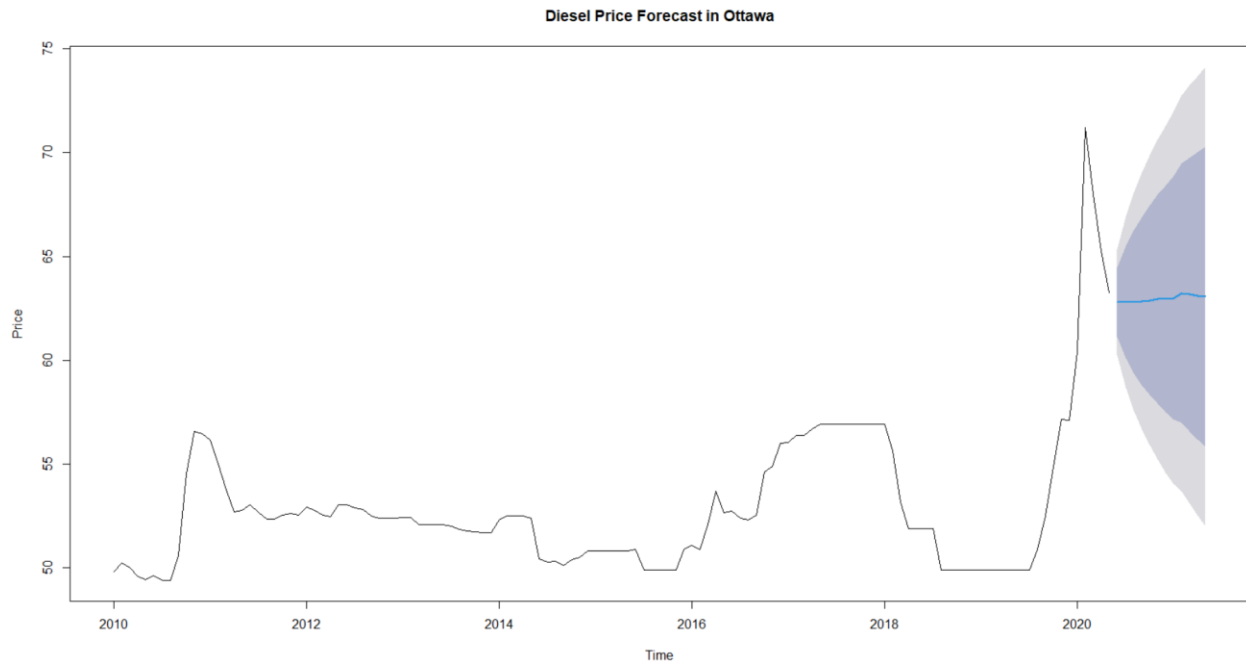
```
forecast (fitARIMA_ts_DSL_Ottawa, h=12)
```

```
> forecast (fitARIMA_ts_DSL_Ottawa, h=12)
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
Jun 2020	62.78422	61.12725	64.44118	60.25011	65.31833
Jul 2020	62.78422	60.09185	65.47659	58.66660	66.90184
Aug 2020	62.80507	59.37693	66.23321	57.56218	68.04795
Sep 2020	62.83842	58.80662	66.87023	56.67231	69.00454
Oct 2020	62.88846	58.33228	67.44463	55.92038	69.85653
Nov 2020	62.93536	57.90922	67.96150	55.24855	70.62218
Dec 2020	62.93432	57.47855	68.39009	54.59045	71.27819
Jan 2021	63.00312	57.14917	68.85706	54.05028	71.95595
Feb 2021	63.22827	57.00155	69.45498	53.70533	72.75121
Mar 2021	63.16114	56.58274	69.73953	53.10035	73.22193
Apr 2021	63.10422	56.19202	70.01643	52.53292	73.67553
May 2021	63.06274	55.83211	70.29336	52.00445	74.12102

Additionally, I can also visualize the forecasting by plotting the prediction.

```
future_DSL_Ottawa <- forecast (fitARIMA_ts_DSL_Ottawa, h=12)
plot (future_DSL_Ottawa, main="Diesel Price Forecast in Ottawa", ylab="Price",
      xlab="Time")
```



In the result graph, the prediction / forecast is illustrated as a blue line. The light shaded area represents the 95% prediction interval and the dark shaded area represents the 80% prediction interval.

## Initial Result

My first prediction of Diesel price for Ottawa marketplace indicates that from the period of June 2020 and May 2021, the price would be stable with a slight increase, and the price would probably stay between 60 and 70 (prices in cents per litre).

I would think that my first try of the forecasting is working. I will apply the same approach to other fuel types and cities. And I will finally generate a comprehensive result to complete my time series analysis project.

All the R codes that I currently used are included here. I will organize my codes for final version and post on GitHub in future weeks.